



Computer Architecture

Final Project: Matrix Chain Multiplication

Speaker : Billy

Advisor : Prof. An-Yeu Wu

Date : 2025/4/28



Overview

- ❖ Given a sequence of matrices, implement **assembly code** to compute the matrix chain multiplication.
- ❖ Your implementation will be scored based on its performance, which is determined by **the number of cycles and cache size**.

$$\begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} \times \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} \times \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} = \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix}$$

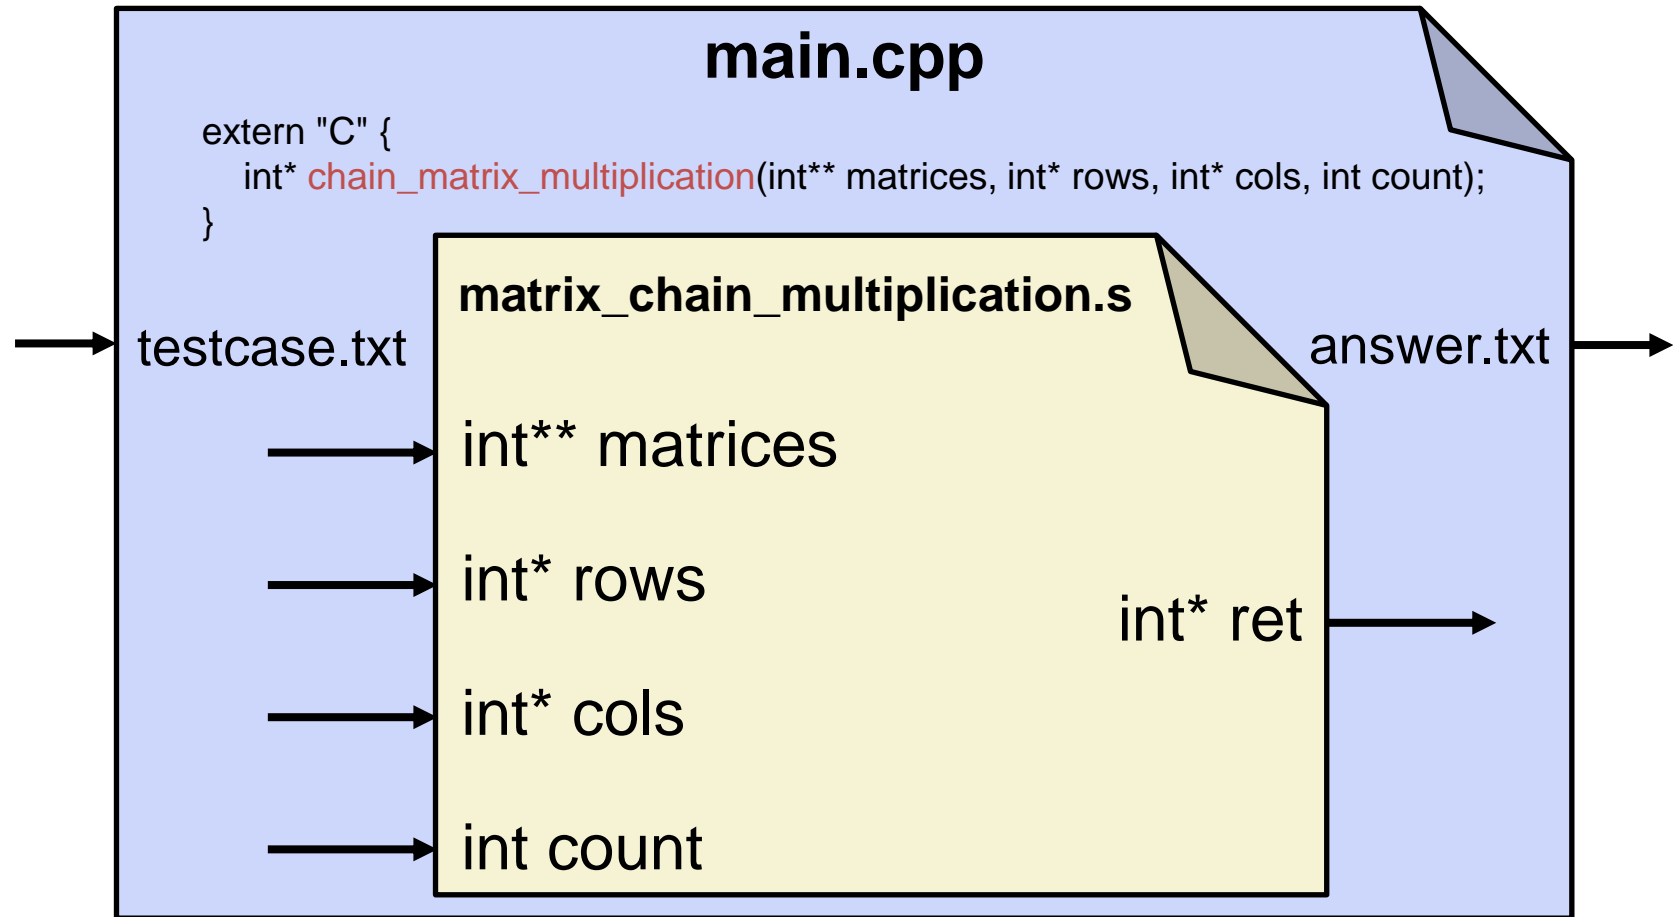


Project Goal

- ❖ Accelerate the function to achieve better performance
 - ❖ Increase hit rate of cache
 - ❖ Use better algorithm (e.g. Dynamic Programming)
- ❖ Performance Formula
 - ❖ $Time \times \left(\log_2 Size_{L1_ICache} + \log_2 Size_{L1_DCache} + \frac{1}{2} \log_2 Size_{L2_Cache} \right)$
 - ❖ Goal: minimize the performance formula



Block Diagram





Input / Output

Type	Port	Reg	I/O	Description
Int **	matrices	x10 (a0)	I	An array storing the addresses of all matrices*.
Int *	rows	x11 (a1)	I	An array storing the row size of each matrix.
Int *	cols	x12 (a2)	I	An array storing the column size of each matrix.
Int	count	x13 (a3)	I	An integer representing the number of matrices.
Int *	ret	x10 (a0)	O	An address storing the result matrix* after computation.

*All matrices are stored in [row-major order](#).

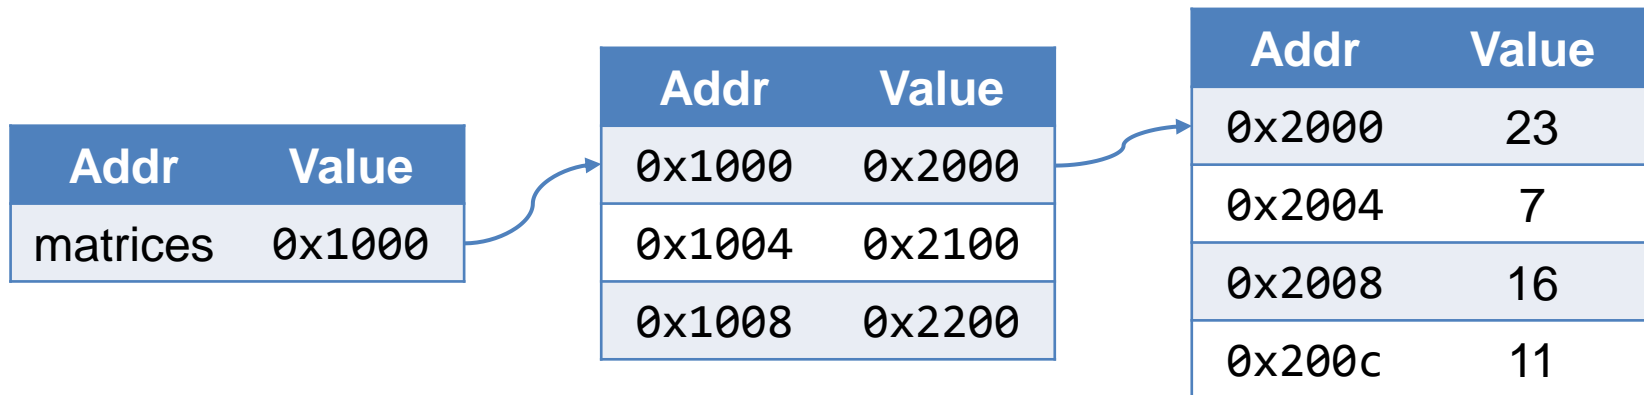
- ❖ In the RISC-V calling convention
 - ❖ Function arguments are passed starting from the a0 register
 - ❖ Return value is stored in the a0 register too



Input

❖ Example:

$$\begin{array}{ccc}
 \begin{bmatrix} 23 & 7 \\ 16 & 11 \end{bmatrix} & \times & \begin{bmatrix} 2 & 13 & 5 \\ 8 & 26 & 18 \end{bmatrix} & \times & \begin{bmatrix} 9 & 15 & 17 \\ 31 & 28 & 4 \\ 14 & 27 & 21 \end{bmatrix} \\
 \nearrow & & \uparrow & & \nwarrow \\
 \text{2x2 matrix} & & \text{2x3 matrix} & & \text{3x3 matrix}
 \end{array}$$





Input

❖ Example:

$$\begin{array}{ccc}
 \begin{bmatrix} 23 & 7 \\ 16 & 11 \end{bmatrix} & \times & \begin{bmatrix} 2 & 13 & 5 \\ 8 & 26 & 18 \end{bmatrix} & \times & \begin{bmatrix} 9 & 15 & 17 \\ 31 & 28 & 4 \\ 14 & 27 & 21 \end{bmatrix} \\
 \nearrow & & \uparrow & & \nwarrow \\
 \text{2x2 matrix} & & \text{2x3 matrix} & & \text{3x3 matrix}
 \end{array}$$

Addr	Value
rows	0x3000

Addr	Value
0x3000	2
0x3004	2
0x3008	3

Addr	Value
cols	0x4000

Addr	Value
0x4000	2
0x4004	3
0x4008	3


Addr	Value
count	3



Output

❖ Example:

$$\begin{bmatrix} 23 & 7 \\ 16 & 11 \end{bmatrix} \times \begin{bmatrix} 2 & 13 & 5 \\ 8 & 26 & 18 \end{bmatrix} \times \begin{bmatrix} 9 & 15 & 17 \\ 31 & 28 & 4 \\ 14 & 27 & 21 \end{bmatrix} = \begin{bmatrix} 19203 & 21505 & 8719 \\ 20286 & 23138 & 9854 \end{bmatrix}$$


 result

Addr	Value
ret	0x9000

Addr	Value
0x9000	19203
0x9004	21505
0x9008	8719
0x900c	20286
0x9010	23138
0x9014	9854

You need to store this value
in a0 before return



RISC-V Calling Convention

- ❖ Input / output start from a0
- ❖ Return address is store in ra
- ❖ **s0 needs to be saved and restored across function call**

Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5–7	t0–2	Temporaries	Caller
x8	s0/fp	Saved register/frame pointer	Callee
x9	s1	Saved register	Callee
x10–11	a0–1	Function arguments/return values	Caller
x12–17	a2–7	Function arguments	Caller
x18–27	s2–11	Saved registers	Callee
x28–31	t3–6	Temporaries	Caller



Memory Allocation

- ❖ Allocation of memory for storing matrices is needed in this project
 - ❖ At least for the result matrix of matrix chain multiplication
- ❖ You can allocate memory by `call malloc` as example below:
 - ❖ Allocate 10 bytes memory space in this example

```
li a0, 10      # Store required memory size in a0 (function argument)
call malloc    # Function call to allocate memory

sw t0, 0(a0)   # Allocated memory address will be stored in a0
               # you can then store data in the allocated memory
```

- ❖ Value in register a0, a1, a2, a3 will be overwritten after this function call
- ❖ Release memory space by `call free`



Matrix Multiplication Assembly Example

- ❖ Basic idea of matrix multiplication: Three loops

```
for (int i = 0; i < n; i++) {  
    for (int j = 0; j < n; j++) {  
        for (int k = 0; k < n; k++) {  
            C[i][j] += A[i][k] * B[k][j]  
        }  
    }  
}
```



Matrix Multiplication Assembly Example

```
.data
A: .word 1, 2, 3, 4          # Matrix A
B: .word 5, 6, 7, 8          # Matrix B
C: .space 16                 # Space for Matrix C (2 x 2 = 4 integers = 16 bytes)

.text
.globl main
main:
    # load base addresses
    la s0, A                  # s0 = base address of A
    la s1, B                  # s1 = base address of B
    la s2, C                  # s2 = base address of C

    # i initialization
    li t0, 0                  # i = 0 (row index)

outer_loop_i:
    # j initialization

outer_loop_j:
    # k initialization
```

Data definition

Address of matrix data

Loop initialization



Matrix Multiplication Assembly Example

```

inner_loop_k:
    li t3, 2                # matrix size = 2

    # Compute address of A[i][k]
    mul t4, t0, t3          # t4 = i * matrix size
    add t4, t4, t2          # t4 = i * matrix size + k
    slli t4, t4, 2          # offset = (i * matrix size + k) * 4
    add t5, s0, t4          # (base address of A) + offset
    lw t6, 0(t5)            # load A[i][k] from memory

    # Compute address of B[k][j]
    # Similar to load A

    # multiplication and accumulation
    mul s7, t6, s6           # A[i][k] * B[k][j]
    add s3, s3, s7           # sum += A[i][k] * B[k][j]

    # Update k = k + 1
    addi t2, t2, 1          # k++
    li s8, 2                # matrix size = 2
    blt t2, s8, inner_loop_k # if k < 2, continue inner loop

    # Compute address of C[i][j]
    # Similar to load A and B

    # Update j = j + 1

    # Update i = i + 1

    # End of program
    jr ra
  
```

} Get $A[i][k]$ address and load

} Get $B[k][j]$ address and load

} Compute $C[i][j]$, iterate k

} Get $C[i][j]$ address and store

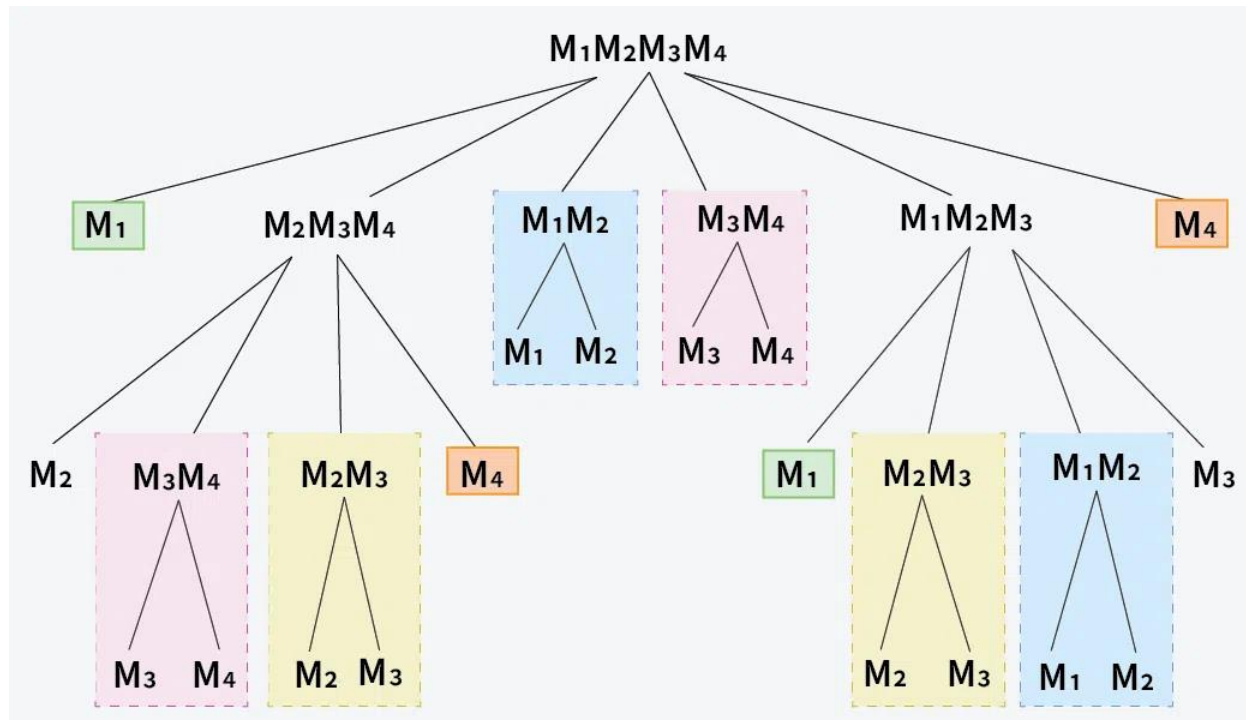
} Iterate i, j

} End of program



Dynamic Programming

- ❖ When multiplying a $k \times m$ matrix with an $m \times n$ matrix
 - ❖ $k \times m \times n$ multiply-add operations are required
 - ❖ Dynamic programming to find the optimal multiplication order




source: www.geeksforgeeks.org



Tiled Matrix Multiplication

- ❖ Tiling matrix into small block that can fit in DCache
 - ❖ Outer loop: iterate through each block of matrices
 - ❖ Inner loop: iterate through elements inside the block
- ❖ Co-optimization of algorithm and cache size



Source: https://www.researchgate.net/figure/Performance-critical-A-B-part-of-the-GEMM-using-a-tiling-strategy-A-thread-iterates_fig1_320499173



Provided Files

Files	Description
Files you should not modify	
main.cpp	Main program
final_config.py	Gem5 configuration file
testbench.py	Python file for checking the correctness
score.py	Python file for calculate performance score
testcase/public/testcase_xx.txt	Public testing data
golden/public/golden_xx.txt	Public golden data
Makefile	Including make commands used in final
Files you need to modify	
matrix_chain_multiplication.s	Main design file
gem5_args.conf	Gem5 argument file, including cache setting



About final_config.py

- ❖ Currently doesn't including config for L2 cache
 - ❖ Implementation of L2 cache is part of HW3 (upcoming)
- ❖ After HW3 due, final_config.py with L2 cache config will be released
 - ❖ Please be aware of the upcoming updates to the final project file
 - ❖ Your final result should be simulated by latest version of final_config.py



Design & Simulation Flow (1)

- ❖ Start docker

```
docker start -i <docker_name>
```

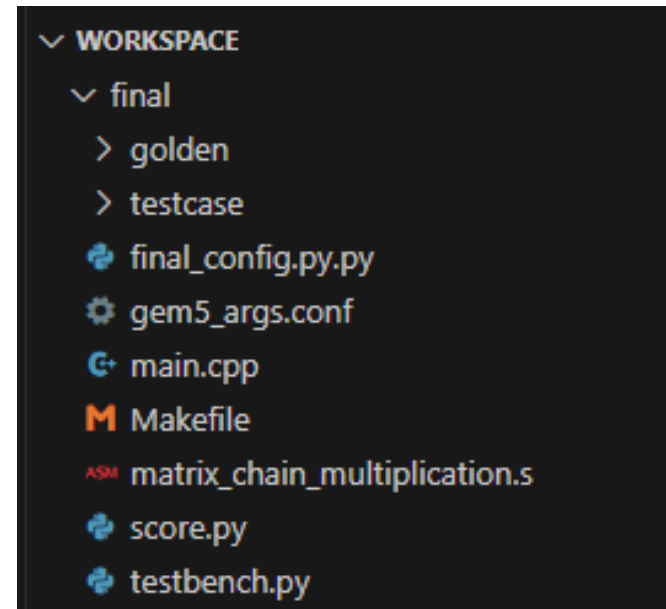
- ❖ Create folder for final

```
mkdir workspace/final
```

```
cd workspace/final
```

- ❖ Put all provided file in workspace/final

- ❖ Complete the design in matrix_chain_multiplication.s





Design & Simulation Flow (2)

- ❖ Modify the cache setting in gem5_args.conf
 - ❖ Settings you can change (mark in red):
 - ❖ GEM5_ARGS = --l1i_size 4kB --l1i_assoc 2 --l1d_size 4kB --l1d_assoc 2 --l2_size 16kB --l2_assoc 4

```
# Modify the GEM5 arguments for the simulation
GEM5_ARGS = --l1i_size 4kB --l1i_assoc 2 --l1d_size 4kB --l1d_assoc 2 --l2_size 16kB --l2_assoc 4
```

- ❖ Compile main.cpp with matrix_chain_multiplication.s

```
make g++_final
```



Design & Simulation Flow (3)

- ❖ Run simulation with all testcases or single testcase

- ❖ This step might take a few minutes

```
make gem5_public_all
```

```
make gem5_public ARGS=P0/P1/P2/P3/P4/P5
```

- ❖ Checking if generated answer.txt match golden.txt

```
make testbench_public
```

- ❖ Compute performance score based on performance formula

- ❖ Read content from m5out/config.json and m5out/out_exec.txt

```
make score_public
```



Test Cases Information

❖ Constraints:

- ❖ Maximum number of matrices (N): 16
- ❖ Maximum dimension of matrices (D): 64
- ❖ Maximum value of matrices (V): 128

❖ Public Test Cases:

- ❖ 00: $N = 3$, $D = 3$, $V = 10$
- ❖ 01: $N = 3$, $D = 5$, $V = 10$
- ❖ 02: $N = 4$, $D = 6$, $V = 16$
- ❖ 03: $N = 5$, $D = 10$, $V = 20$
- ❖ 04: $N = 6$, $D = 8$, $V = 12$
- ❖ 05: public performance test



Group Formation

- ❖ 2 students per group
- ❖ Find your groupmate and submit the form before **5/4 (Sun.) 23:59**
 - ❖ <https://forms.gle/rKvznPAk3oMZWjWdA>
- ❖ You may use NTU COOL 討論區 to find your groupmate
 - ❖ https://cool.ntu.edu.tw/courses/45288/discussion_topics/393734
- ❖ For those who haven't submit the form, the TAs will randomly assign groupmate for you
- ❖ Email to andrew@access.ee.ntu.edu.tw for any private questions related to group formation
 - ❖ Subject should start with [113-2 CA 分組]



Submission

- ❖ Deadline: **2025/6/8 23:59:59 (UTC+8)**
- ❖ matrix_chain_multiplication.s
 - ❖ Your assembly code that implement matrix chain multiplication
- ❖ gem5_args.conf
 - ❖ Your simulation arguments (cache size & cache associative)
- ❖ report.pdf
 - ❖ 1-2 pages description of your design spec, special techniques used, team's division of work, and reflections.
- ❖ Submit 組內互評 form
 - ❖ Contribution score for you and your teammate



Submission – Group



- ❖ Compress teamID_final/ in a zip file
 - ❖ named teamID_final_vk.zip (k: version number, e.g., 1,2,...)
- ❖ Upload teamID_final_vk.zip to NTUCOOL (ex: team01_final_v1.zip)

- ❖ File structure:

- ❖ teamID_final_vk.zip

- ❖ teamID_final/
 - matrix_chain_multiplication.s
 - gem5_args.conf
 - report.pdf

Example:



- ❖ Incorrect file name or format would get **10% penalty**



Submission – Individual

- ❖ Submit 組內互評 form
 - ❖ <https://forms.gle/x2cKjfPaR9GKfmji8>
- ❖ Contribution score
 - ❖ 0 = no contribution
 - ❖ 10 = full contribution
 - ❖ Summation of all the group should be 10
- ❖ **Everyone** needs to submit this form !



Grading Policy

- ❖ Baseline 40% + Performance 40% + Report 20%
- ❖ 0 points for late submission or plagiarism
- ❖ -10% for incorrect naming or format

Item	%	Description
Public Testcases	30	Pass all 6 public testcases (5% each)
Private Testcases	10	Pass all 5 private testcases (2% each)
Public Ranking	20	Assign scores based on the PR value of your performance score of P5 (public).
Private Ranking	20	Assign scores based on the PR value of your performance score of P10 (private).
Report	15	1~2 pages
Contribution	5	Based on 組內互評



Discussion

- ❖ NTU COOL Discussion Forum: Final project 討論區
 - ❖ TAs will prioritize answering questions on the NTU COOL discussion forum

- ❖ Email: r13943002@ntu.edu.tw
 - ❖ Title should start with [113-2 CA Final Project]



Reference

- ❖ Matrix-chain multiplication
 - ❖ [Matrix Chain Multiplication | GeeksforGeeks](#)
- ❖ Dynamic Programming
 - ❖ [Dynamic Programming or DP | GeeksforGeeks](#)
- ❖ Tiled matrix multiplication
 - ❖ [penny-xu.github.io/blog/tiled-matrix-multiplication/](#)